

Edges & Features

Computer Vision

From Edges to Boundaries

edge: image location of rapid change in intensity
(caused by discontinuities of surface, depth, or illumination)

edges encode most shape and semantic information


to detect edge pixels: image derivatives (see sharpening filters)

follow-up steps:

→ edge linking (connect edge pixels to edge segments)

→ boundary detection

→ non-semantic image segmentation and object recognition




Canny Edge Detector

Noise Complications for Edge Detection

image derivatives accentuate high frequencies (high-pass filter)


→ amplification of noise (proportion of noise to signal larger at high frequencies)




Noisy input image

intensity of pixels

$$f(x)$$




$$\frac{d}{dx}f(x)$$



Need for Band-Pass Filter



Image with Edge



Edge Location

idea:

blurring (low-pass filter) before
derivation (high-pass filter)






Image + Noise



Derivatives detect
edge *and* noise




Smoothed derivative removes
noise, but blurs edge

edge detector should be isotropic:
need for circularly symmetric
smoothing filter


→ Gaussian is only separable one

Gradient with Gaussian Smoothing



same result, but more efficient:

$$\frac{d}{dx}(f * G) = f * \frac{dG}{dx}$$



to find edges, look for peaks
in first-order derivatives


Edge Strength and Orientation

gradient magnitude (edge strength):

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

gradient direction (perpendicular to edge orientation):

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$




Non-Maximum Suppression

only detect as edge pixel if local maximum along gradient direction (using interpolation between pixels)

→ thin the edges

(can be done with Laplacian zero-crossings)



gradient magnitude:

before



after

Thresholding

still need to decide for each pixel if edge pixel or not

→ generate binary image by thresholding gradient magnitudes

problem: thresholding can lead to many disconnected edges

also, choice of threshold value can be difficult:

high threshold:
few edges



low threshold:
many edges

Hysteresis: Edge Tracking


define two thresholds t_1 and t_2 :

$\|\nabla f\| < t_1$ no edge pixel

$\|\nabla f\| > t_2$ strong edge
→ edge pixel


$t_1 < \|\nabla f\| < t_2$ weak edge
→ uncertain

for weak edges, decide on neighborhood:



Putting It Together: Canny Edge Detector

- smoothing by Gaussian filter
- calculate magnitude and direction of gradient image
- non-maximum suppression (reduce to single-pixel edges)
- hysteresis (binary edge image with linked edge pixels)



Scale Space

besides threshold values, one more important parameter in Canny edge detector: used standard deviation of Gaussian smoothing filter

- broad Gaussians detects large-scale edges
 - narrow Gaussians detect fine edges
- can be combined to detect and select edges at different scales



original

Canny with $\sigma = 1$

Canny with $\sigma = 2$





Marr-Hildreth Edge Detector

Alternative: Laplacian of Gaussian (LoG)

use same idea of smoothing by Gaussian filter before Laplacian

$$\nabla^2(f * G) = f * \nabla^2 G$$

Laplacian:
zero-crossings at edges



0	0	-1	0	0
0	-1	-2	-1	0
-1	-2	16	-2	-1
0	-1	-2	-1	0
0	0	-1	0	0

approximated
LoG filter


scale of edge detection defined by σ

Zero-Crossings for Edge Detection

use zero-crossings for edge detection
(instead of maxima for gradient)

effect of Gaussian smoothing:

LoG



Laplacian



zero-crossings also thresholded:



Image (I)



Laplacian
(0 maps to 128)



Laplacian
"Zero Crossings"

LoG vs Derivative of Gaussian



LoG




derivative of Gaussian



Derivative of Gaussian (∇G)



Laplacian of Gaussian ($\nabla^2 G$)



$$\frac{\partial^2}{\partial x^2}(n_\sigma) + \frac{\partial^2}{\partial y^2}(n_\sigma)$$

Inverted "Sombrero"
(Mexican Hat)

- Laplacian provides only edge location, no strength or direction
- but requires only one convolution instead of two

Difference of Gaussians (DoG)


difference of two Gaussians with different σ as approximation to LoG

→ less complex and faster

ratio here 1.6:1

solid: LoG
dashed: DoG

Marr-Hildreth edge detector




Comparison of Edge Detectors

original




thresholded gradient of
smoothed image



Marr-Hildreth
edge detector



Canny edge detector



Hough Transform

Boundary Detection

In fact, this uses
semantic segmentation:



one option for image segmentation: detect boundaries

but edges are no boundaries

simplest approach:

curve fitting across edge pixels

in general, only possible for local regions (requiring pre-knowledge),
as globally checking all possible lines is computationally prohibitive



Idea: Curves of Specified Shape

Lines:

variables

$$y = m \cdot x + b$$

parameters

variables

$$b = -x \cdot m + y$$

parameters

only edge pixels:





Image space



Parameter space

one line in
image space

e.g.

→ choose parameter values m and b (binned) with maximum line crossings in parameter space

Hough Transform: Edges to Lines

issue: parameter space of m and b infinite

→ use normal representation of lines: $\rho = x \cdot \cos \theta + y \cdot \sin \theta$






Image space




Parameter space

accumulator cells:



Examples





line

rectangle

circle

application:




lines corresponding
to accumulator cells
with most entries



Circle Detection

If radius r is known: Accumulator Array: $A(a, b)$

Image Space




Parameter Space




If radius r is NOT known: Accumulator Array: $A(a, b, r)$

Image Space




Parameter Space



Feature-Based Object Recognition

Hough transform can be used to detect shape features

more complex (predefined) shapes via generalized Hough transform



Hough detection


but struggling with cluttered scenes or different viewpoints

Non-Semantic Segmentation

image segmentation: clustering of pixels

various approaches:

- detect boundaries between regions based on discontinuities in intensity levels (edges)
- active contours
- thresholding using intensity histograms:
- region splitting or merging
- graph-based segmentation
- ...




Local Features

Keypoint Detection & Matching


different stages:

- **feature detection:** search for distinguishing keypoints (aka interest points) in images
- **feature description:** construct (invariant) descriptors/features from regions (pixel patches) around detected keypoint locations
- **feature matching:** search for matching keypoints (correspondences) in other images



Why Local Keypoint Features for Instance Recognition?

often non-obvious shapes




NASA Mars Rover images with SIFT feature matches

at different scales, under perspective changes and with partial observability

→ matching predefined shapes or large patches very difficult

Good and Bad Keypoints



flat regions cannot
be localized


edges are difficult
to localize



corners are interesting
candidates


Feature Detection

Harris Corner Detector




1. define a small window and shift over image
2. for each region calculate image gradients

discrete convolution with weighting kernel w :



or better: Gaussian
→ rotation invariant



$$w * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

directions of maximum variance


Flat Region




Edge Region




Corner Region



$\lambda_1 \sim \lambda_2$
Both are Small




$\lambda_1 \gg \lambda_2$
 λ_1 is Large
 λ_2 is Small




$\lambda_1 \sim \lambda_2$
Both are Large

3. build covariance matrix of image gradients and calculate its eigenvalues λ_1 and λ_2 (aka **Principal Component Analysis**)



response function R




alternative:
 use the smaller eigenvalue
 as response function

$$R = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2$$


where: $0.04 \leq k \leq 0.06$

4. use threshold on eigenvalues for corner detection
5. perform non-maximum suppression


response function:



thresholded response:




corner response rotation invariant



eigenvalues stay the same

but not scale invariant



corner!





important for detection in images
with different sizes or at positions
with different depths

Scale-Invariant Keypoint Detection

→ need to detect keypoints at the right scale for each image


→ search for both position and scale






Additional Keypoint-Detection Requirements

corners often not descriptive enough for complex objects (e.g., smoother regions)

→ need to detect not only corners, but all kinds of **blobs** (keypoint as center of blob region)



and it must be possible to find the right scale for the blob to be detected:



LoG Operator

Laplacian

$$\nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$$

Gaussian



n_σ

LoG




$\nabla^2 n_\sigma$


NLoG



$\sigma^2 \nabla^2 n_\sigma$



Detect Blobs at Different Sizes



example: LoG filter for $\sigma = 1.4$

0	0	1	2	2	2	1	0	0
0	1	3	5	5	5	3	1	0
1	3	5	3	0	3	5	3	1
2	5	3	-12	-23	-12	3	5	2
2	5	0	-23	-40	-23	0	5	2
2	5	3	-12	-23	-12	3	5	2
1	3	5	3	0	3	5	3	1
0	1	3	5	5	5	3	1	0
0	0	1	2	2	2	1	0	0

need larger filter for larger σ

→ need to compute this at many scales

Full size




3/4 size




multiple scales via Gaussian
filters with different σ

 σ


2.1




4.2




6.0




9.8



15.5




17.0




Full size image


2.1




4.2



6.0




9.8



15.5




17.0




3/4 size image

DoG as Approximation of LoG





Scale Invariant Feature Transform (SIFT) Detector


look at $3 \times 3 \times 3$
neighborhoods




detect as keypoint if
center is extrema in
this pixel-scale volume




detected keypoints
with different scales



SIFT
Interest Points



Gaussian Pyramid for Different Image Sizes



Feature Description & Matching

Motivation

feature descriptor:

simplified representation of an image containing only its most important information

raw pixel intensities difficult to interpret

(→ deep learning)

Local Patches as Features

extract patches around keypoints

→ descriptor: vector summarizing content of these patches

simplest descriptor: just raw intensity values of pixels in patch
(template matching)

but need to be invariant to translation, rotation, scale, (illumination)

SIFT Descriptor

16 × 16 square window around
detected keypoint
(8 × 8 shown here for visualization)

Keypoint neighborhood





Image gradients


weighted by Gaussian
fall-off function

Array of Orientation
Histograms




Keypoint descriptor

rotation invariance:
rotate gradients relative to the
keypoint's principal orientation before
calculating the descriptor



Keypoint neighborhood



SIFT descriptor is vector of length 128:


- 4 × 4 histograms per keypoint
- 8 orientation bins per histogram

using the image of the detected scale (i.e., smoothed
with corresponding Gaussian filter) for this keypoint


Many Other Feature Detectors/Descriptors

- MOPS: simple normalized intensity patches
- SURF, ORB: similar to SIFT, but faster
- HOG: used for larger image regions, particularly suited for human detection
- ...

Multi-scale Oriented Patches:



Histogram of Oriented Gradients:



counting occurrences of
gradient orientations in
localized portions of an image

features covering entire image
(instead of just local patches)

Feature Matching

goal: for each keypoint in a given image,
find nearest neighbor in other image
(or rank different potential matches)

→ look for smallest distance d between
the corresponding feature descriptors $\mathbf{f}$
(vectors) in image 1 and 2

several possible distance metrics, e.g.,
L2:

$$d = \sqrt{(\mathbf{f}_1 - \mathbf{f}_2)^2}$$





image alignment: homography ...

Object Recognition

two database
images to be
recognized:




using SIFT features



approach: extract features and compare against stored object features in database (successful if defined number of feature matches found), subsequent match verification by geometric alignment

works well for instance recognition (e.g., specific car), but class recognition (e.g., cars in general) requires more abstraction → ML



Facial Recognition with Eigenfaces

alternative approach for instance recognition: eigenimages

most popular example: individual face images expressed as combination of basis images (called eigenfaces)

PCA (unsupervised ML) on large training set of face images, use principal components (eigenvectors with largest eigenvalues) as basis images

identification: compare vectors of eigenface components (projections to eigenspace), typically using ~100 eigenfaces

some eigenfaces:

